

# Database Management System (DBMS)

## Lecture-42

---

Dharmendra Kumar

December 28, 2020

## Concurrency Control

---

# Concurrency Control

---

When several transactions execute concurrently in the database, however, the isolation property may no longer be preserved. To ensure that it is, the system must control the interaction among the concurrent transactions. The mechanism used to control the interaction of transactions is called concurrency control scheme.

There are number of concurrency control schemes.

1. Lock based protocol
2. Time stamp based protocol
3. Validation based protocol
4. Multiple granularity protocol
5. Multi-version protocol

## Lock based protocol

---

A lock is a mechanism to control concurrent access to a data item. Data item can be locked in two modes.

**Shared mode(S):** If a transaction  $T_i$  has obtained a shared-mode lock on item Q, then  $T_i$  can read, but cannot write, Q.

**Exclusive mode(S):** If a transaction  $T_i$  has obtained an exclusive-mode lock on item Q, then  $T_i$  can both read and write Q.

**Note:** The transaction makes the lock request to the concurrency control manager. Transaction can process only after lock request is granted.

# Concurrency Control

## Compatibility function

Given a set of lock modes, we can define a compatibility function on them as follows.

Let A and B represent arbitrary lock modes. Suppose that a transaction  $T_i$  requests a lock of mode A on item Q on which transaction  $T_j$  ( $T_i \neq T_j$ ) currently holds a lock of mode B. If transaction  $T_i$  can be granted a lock on Q immediately, in spite of the presence of the mode B lock, then we say mode A is compatible with mode B. Such a function can be represented conveniently by a matrix. An element  $\text{comp}(A, B)$  of the matrix has the value **true** if and only if mode A is compatible with mode B.

Compatibility Matrix

	S	X
S	true	false
X	false	false

# Concurrency Control

## Note:

- A transaction requests a shared lock on data item Q by executing the lock-S(Q) instruction.
- A transaction requests an exclusive lock on data item Q by executing the lock-X(Q) instruction.
- To unlock the data item Q, we use unlock(Q) instruction.

## Note:

To access a data item, transaction  $T_i$  must first lock that item. If the data item is already locked by another transaction in an incompatible mode, the concurrency control manager will not grant the lock until all incompatible locks held by other transactions have been released. Thus,  $T_i$  is made to wait until all incompatible locks held by other transactions have been released.

# Concurrency Control

**Example:** Consider the following two transactions  $T_1$  and  $T_2$  with locking modes.

## Transaction $T_1$ and $T_2$

```
 $T_1$ : lock-X( $B$ );  
      read( $B$ );  
       $B := B - 50$ ;  
      write( $B$ );  
      unlock( $B$ );  
      lock-X( $A$ );  
      read( $A$ );  
       $A := A + 50$ ;  
      write( $A$ );  
      unlock( $A$ ).
```

```
 $T_2$ : lock-S( $A$ );  
      read( $A$ );  
      unlock( $A$ );  
      lock-S( $B$ );  
      read( $B$ );  
      unlock( $B$ );  
      display( $A + B$ ).
```

# Concurrency Control

Consider the following schedule-1 of these transactions.

## Schedule-1

$T_1$	$T_2$	concurrency-control manager
lock-X(B)		grant-X(B, $T_1$ )
read(B)		
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, $T_2$ )
	unlock(A)	
	lock-S(B)	
	read(B)	grant-S(B, $T_2$ )
	unlock(B)	
	display(A + B)	
lock-X(A)		grant-X(A, $T_2$ )
read(A)		
$A := A + 50$		
write(A)		
unlock(A)		



## Concurrency Control

Suppose that the values of accounts A and B are \$100 and \$200, respectively. If these two transactions are executed serially, either in the order  $T_1, T_2$  or the order  $T_2, T_1$ , then transaction  $T_2$  will display the value \$300. If, however, these transactions are executed concurrently, then schedule 1 is possible. In this case, transaction  $T_2$  displays \$250, which is incorrect. The reason for this mistake is that the transaction  $T_1$  unlocked data item B too early, as a result of which  $T_2$  saw an inconsistent state.

# Concurrency Control

**Example:** Consider the following two transactions  $T_3$  and  $T_4$  with locking modes.

## Transaction $T_3$ and $T_4$

```
 $T_3$ : lock-X( $B$ );  
      read( $B$ );  
       $B := B - 50$ ;  
      write( $B$ );  
      lock-X( $A$ );  
      read( $A$ );  
       $A := A + 50$ ;  
      write( $A$ );  
      unlock( $B$ );  
      unlock( $A$ ).
```

```
 $T_4$ : lock-S( $A$ );  
      read( $A$ );  
      lock-S( $B$ );  
      read( $B$ );  
      display( $A + B$ );  
      unlock( $A$ );  
      unlock( $B$ ).
```

# Concurrency Control

Consider the following schedule-2 of these transactions.

Schedule-2

$T_3$	$T_4$
lock-X(B)	
read(B)	
$B := B - 50$	
write(B)	
	lock-S(A)
	read(A)
	lock-S(B)
lock-X(A)	

Consider the partial schedule-2 for  $T_3$  and  $T_4$ . Since  $T_3$  is holding an exclusive-mode lock on B and  $T_4$  is requesting a shared-mode lock on B,  $T_4$  is waiting for  $T_3$  to unlock B. Similarly, since  $T_4$  is holding a shared-mode lock on A and  $T_3$  is requesting an exclusive-mode lock on A,  $T_3$  is waiting for  $T_4$  to unlock A. Thus, we have arrived at a state where neither of these transactions can ever proceed with its normal execution. This situation is called **deadlock**.

**Note:** When deadlock occurs, the system must roll back one of the two transactions.

**Locking protocol** This is the set of rules indicating when a transaction may lock and unlock each of the data items.

**Note:** A schedule  $S$  is legal under a given locking protocol if  $S$  is a possible schedule for a set of transactions that follow the rules of the locking protocol.

**Note:** A locking protocol ensures conflict serializability if and only if all legal schedules are conflict serializable.

## Starvation

Suppose a transaction  $T_2$  has a shared-mode lock on a data item, and another transaction  $T_1$  requests an exclusive-mode lock on the data item. Clearly,  $T_1$  has to wait for  $T_2$  to release the shared-mode lock. Meanwhile, a transaction  $T_3$  may request a shared-mode lock on the same data item. The lock request is compatible with the lock granted to  $T_2$ , so  $T_3$  may be granted the shared-mode lock. At this point  $T_2$  may release the lock, but still  $T_1$  has to wait for  $T_3$  to finish. But again, there may be a new transaction  $T_4$  that requests a shared-mode lock on the same data item, and is granted the lock before  $T_3$  releases it. In fact, it is possible that there is a sequence of transactions that each requests a shared-mode lock on the data item, and each transaction releases the lock a short while after it is granted, but  $T_1$  never gets the exclusive-mode lock on the data item. The transaction  $T_1$  may never make progress, and is said to be starved. This situation is said to be **starvation**.

## Two-phase locking protocol

---

This protocol requires that each transaction issue lock and unlock requests in two phases:

- 1. Growing phase:** A transaction may obtain locks, but may not release any lock.
- 2. Shrinking phase:** A transaction may release locks, but may not obtain any new locks.

Initially, a transaction is in the growing phase. The transaction acquires locks as needed. Once the transaction releases a lock, it enters the shrinking phase, and it can issue no more lock requests.

## Concurrency Control

**Example:** Transactions  $T_3$  and  $T_4$  are locked in two phase. While, transactions  $T_1$  and  $T_2$  are not locked in two phase.

**Note:** Two-phase locking protocol ensures conflict serializability. The serializability order of transactions will be based on lock point in the transactions.

**Lock point:** Lock point of a transaction is a point in the schedule where the transaction has obtained its final lock (the end of its growing phase).

**Note:** Two-phase locking does not ensure freedom from deadlock. Observe that transactions  $T_3$  and  $T_4$  are in two phase, but, in schedule 2, they are deadlocked.

**Note:** In addition to being serializable, schedules should be cascadeless. Cascading rollback may occur under two-phase locking.

## Concurrency Control

**Example:** Consider the partial schedule in the following figure:-

### Partial schedule under two-phase locking protocol

$T_5$	$T_6$	$T_7$
lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A)	lock-X(A) <b>read(A)</b> write(A) unlock(A)	lock-S(A) read(A)



## **Strict two-phase locking protocol**

This protocol requires not only that locking be two phase, but also that all exclusive-mode locks taken by a transaction be held until that transaction commits.

## **Rigorous two-phase locking protocol**

Another variant of two-phase locking is the rigorous two-phase locking protocol, which requires that all locks be held until the transaction commits.

**Note:** With rigorous two-phase locking, transactions can be serialized in the order in which they commit.

## Lock Conversion

**Upgrade:** We denote conversion from shared to exclusive modes by upgrade.

**Downgrade:** We denote conversion from exclusive to shared by downgrade.

**Note:** Lock conversion cannot be allowed arbitrarily. Rather, upgrading can take place in only the growing phase, whereas downgrading can take place in only the shrinking phase.

**Note:** Strict two-phase locking and rigorous two-phase locking (with lock conversions) are used extensively in commercial database systems.

## Concurrency Control

**Note:** A simple but widely used scheme automatically generates the appropriate lock and unlock instructions for a transaction, on the basis of read and write requests from the transaction:

- When a transaction  $T_i$  issues a read(Q) operation, the system issues a lock- S(Q) instruction followed by the read(Q) instruction.
- When  $T_i$  issues a write(Q) operation, the system checks to see whether  $T_i$  already holds a shared lock on Q. If it does, then the system issues an upgrade( Q) instruction, followed by the write(Q) instruction. Otherwise, the system issues a lock-X(Q) instruction, followed by the write(Q) instruction.
- All locks obtained by a transaction are unlocked after that transaction commits or aborts.